

DocXAR-CUFIT & UniMoE-PEFT: Robust Parameter-Efficient Fine-Tuning for Multimodal Foundation Models

MD. Rafi Imam | Software Development Officer (AI & Automation), Jamuna Bank PLC

Email: rafiimam96@gmail.com | Portfolio: rafiimam.github.io/Personal-website/ | B.Sc. ECE, RUET (3.55 Last-60 GPA) | GRE Quant: 164 | IELTS: 7.0

SUMMARY: This report outlines the development and validation of the **DocXAR-CUFIT** framework, designed to mitigate structured OCR noise in Multimodal Document Understanding (MDU). By integrating Layer-Wise Sensitivity Masking, Cross-Attention Adapters, and a Decoupled Quality Estimator, the framework recovers significant performance degradation caused by noisy inputs. Multi-seed evaluations on the FUNSD dataset under 50% synthetic OCR corruption demonstrate a highly stable Mean Test F1 Score of **0.7992 ± 0.0067**, offering a paradigm-shifting improvement over standard noisy LoRA adaptation (~0.6334).

1. PROBLEM FRAMING AND RESEARCH HYPOTHESIS

The evolution of Multimodal Foundation Models (MFMs) has necessitated a strategic pivot from addressing classical "noisy-label" learning—where ground-truth labels are corrupted—to the more arduous domain of **noisy-input adaptation**. Within Multimodal Document Understanding (MDU), structured OCR noise represents a critical failure point. Unlike stochastic pixel noise, OCR degradation selectively erodes the linguistic coherence of tokens while preserving visual-spatial layout features, creating a modality misalignment that standard Parameter-Efficient Fine-Tuning (PEFT) methods fail to adjudicate.

Standard LoRA updates weights uniformly without regard for the reliability of input modalities, leading to a "garbage-in, garbage-out" propagation across the transformer backbone. **DocXAR-CUFIT** addresses this by combining cross-attention adapters with layer-wise sensitivity masking, enforcing a uniform 25% active parameter density to isolate clean visual-spatial signals and prevent rank collapse under severe text degradation.

2. TECHNICAL METHODOLOGY: REFINEMENTS IN DOCXAR-CUFIT

Maintaining model stability during the fine-tuning of 12-layer transformer blocks requires precise calibration of gradient flow. Without structural constraints, models often succumb to global gradient starvation, where specific layers saturate early. DocXAR-CUFIT introduces three core architectural mechanisms:

- Layer-Wise Sensitivity Masking:** Enforces a uniform 25.0% active parameter density across all transformer blocks via tensor-by-tensor thresholding.
- Cross-Attention Adapter:** Connects question embeddings directly to document key-value pairs after the final encoder block, filtering out non-relevant corrupted spans.
- Decoupled Quality Estimator $Q(x)$:** Evaluates per-sample input corruption under `torch.no_grad()` with batch normalization to prevent weight collapse.




Figure 1: DocXAR-CUFIT Architectural Pipeline detailing input masking, cross-attention adaptation, and quality estimation.

Metric / Parameter	Pre-Registration Status	Post-Registration Refined Status	Architectural Impact
Active LoRA Density	Global (Unconstrained)	25.0% (Uniform Layer-wise)	Prevents Rank Collapse
Residual Scalar (s_i)	Frozen ($s_i = 1.0$)	Active / Unfrozen	Dynamic Layer Reweighting
Avg. LoRA Weight Norm	2.308	3.8517 (+66.9%)	Expands Noise Capacity
Span Selection Strategy	Greedy Selection	Top-K Re-Ranked	Filters OCR Span Errors

Table 1: System Refinements and Parameter Norm Expansion under DocXAR-CUFIT.

3. EXPERIMENTAL RESULTS & MULTI-SEED VALIDATION

Statistical reproducibility is the cornerstone of PEFT research, particularly when noise models introduce high variance across training runs. We conducted multi-seed evaluations on LayoutLMv3 using the FUNSD dataset under 50% synthetic OCR corruption (evaluating Seeds 42, 123, and 2024). DocXAR-CUFIT achieved a highly stable **Mean Test F1 Score of 0.7992 ± 0.0067** (Peak Seed 123 F1: **0.8067**), representing a massive **+26.2% relative improvement** over standard noisy LoRA adaptation (~0.6334).




Figure 2: FUNSD Multi-Seed Benchmark Results comparing Reference, Noisy LoRA, and DocXAR-CUFIT Seed Runs.

4. FUTURE VISION: UNIMOE-PEFT (UNCERTAINTY-GATED MIXTURE-OF-EXPERTS)

The next stage of this research involves transitioning from sample-level scalar reweighting to token-level cross-modal uncertainty routing through **UniMoE-PEFT (Uncertainty-Gated Mixture-of-Experts)**.

4.1 Token-Level Modality Discrepancy Index (TMDI_t)

We are formulating a Token-Level Modality Discrepancy Index (TMDI_t) leveraging cross-attention maps and Gaussian region uncertainty. By modeling visual bounding boxes as a mixture of Gaussians, TMDI_t quantifies fine-grained OCR degradation at the individual token level, identifying exactly where visual text rendering diverges from extracted OCR tokens.

4.2 Noise-Gated MoE Router & PAC-Bayesian Bounds

The TMDI_t metric feeds directly into a noise-gated router that dynamically directs clean tokens to specialized text experts while routing corrupted tokens to visual-spatial experts. Mathematically, this formulation is constrained by PAC-Bayesian generalization bounds of order $O(\sqrt{(k/n)})$, proving that sparse parameter adaptation under input noise maintains bounded generalization error.

Summary Conclusion & Industry Impact: DocXAR-CUFIT proves that parameter-efficient adaptation can recover from severe input corruption without full backbone retraining or clean OCR dependencies. This enables fault-tolerant Document AI deployment in edge and enterprise environments where high-fidelity OCR is computationally prohibitive. Future work on UniMoE-PEFT extends these principles to general multimodal foundation models.